

The Greek Poverty Paradox

Why Greece's official poverty statistics miss part of the story — and why Greek pessimism is not just mood.

15 chapters Companion to the full technical report

Data vintage 2026-08-19 (work-effort series: 2026-08-20)


Two out of three Greek households say they struggle to make ends meet. That is the highest share anywhere in the European Union — more than triple the EU average — and it has been true, without a single exception, every year since 2011.

Ask Greece's official poverty statistics instead, and the picture looks far less extreme. The EU's standard income-poverty rate, AROP, ranks Greece 4th-highest of 27 EU states: elevated, but not alone at the top, not close to it. That gap — between how bad the official number says things are, and how bad Greek households say things feel — is the puzzle this piece exists to take apart.

It would be easy to conclude that one of these numbers must be wrong — exaggerated pessimism on one side, or a blind statistic on the other. That is the wrong instinct. **They are measuring different things.** AROP measures relative income within Greece: how far a household's income sits below the middle of the pack, in whatever year you happen to be counting. Households, when Eurostat asks them directly, are answering a broader question — whether life feels affordable and stable. Nothing requires those answers to move together. This is the story of what happens when they stop, and how much of the resulting gap can actually be accounted for.

47.6pt

The gap between Greece's felt poverty and AROP, the EU's standard income-poverty measure, in 2025 — more than three times the next-largest gap in the EU. Bulgaria, in second place, sits at **14.9 points**.



There is a cruder way to see how far Greece fell, and it needs no statistics at all. Take every measure where it is obvious which direction is the bad one — unemployment, wages, what prices cost against a local paycheck, whether people can heat their homes, what they expect of next year — and ask one question of each year: on how many of them is Greece in the worst fifth of the Union? Before the crisis the answer was about one in five. From 2012 onward it has been closer to three in five, and in 2024 it reached two-thirds. No other member state is near it; Bulgaria, the next highest, sits at half. The striking part is not the level but the shape. The jump happens in 2011 and then simply stays. Fourteen years on, Greece has not climbed back out.

A CAUTION ABOUT THAT NUMBER

It describes the situation; it does not explain it. When the same measure was built as a predictor and tested properly, it failed — it tells you nothing about *why* Greek households report what they report, and this report does not use it that way. Two things were done to keep it from flattering the story: whether "high" or "low" counts as bad was decided for each measure in advance rather than after looking, and every measure the models themselves rely on was thrown out first, along with the subjective-poverty question at the heart of the puzzle. Otherwise the number would just be the puzzle, restated.

The EU already has a partial answer to "isn't AROP too narrow." **AROPE** — a broader official measure adding material deprivation and very-low household work intensity to income poverty — is Eurostat's own recognition that income alone doesn't capture everything. It narrows Greece's gap, from 47.6 points to 39.7 — still the largest gap of any EU country by a wide margin, just a smaller one. AROPE is a bridge in this piece, not a replacement for AROP: it motivates Part II's search for what a broader view of hardship might explain, without this piece claiming to rebuild AROPE itself from the pieces tested along the way. Eurostat doesn't publish how AROPE's three components overlap inside a single household, so nothing here can honestly claim to reconstruct it — only to test the same underlying idea, that income poverty alone is too narrow a lens, one hardship dimension at a time.

What follows traces that idea in order. First, how much of the raw AROP gap is arithmetic — a ruler that shrank at the worst possible moment (Chapter 1) — and how much survives being tested against comparable EU countries, using AROP throughout as the primary measure (Chapter 2). Then what that leftover residual is actually made of: an economy, a paycheck, prices, and a home that never fully recovered (Chapters 3–6), the single sharpest reason why (Chapter 7), and underneath that, how much of it is simply the accumulated weight of carrying all of it for sixteen years instead of three or four (Chapter 8). Then what it did to people beyond their bank accounts, including a national average that turns out to hide a sharp split by age and household type (Chapters 9–11), whether having a

job actually solves it (Chapter 12), what wasn't it (Chapter 13), one more honest stress-test of the whole premise (Chapter 14), before the story lands (Chapter 15).

A ruler that shrank

The official poverty line isn't fixed. It moves with the very economy it's supposed to be measuring.

Here is the quiet flaw at the center of the EU's standard poverty measure: it is not an income level. It is a percentage — 60% of whatever the national median income happens to be, *that year*. In an ordinary economy, where incomes drift up slowly and roughly together, that is a reasonable way to define "poor relative to your neighbors."

Greece's economy, starting in 2010, was not ordinary. Incomes across the whole country fell together, hard and fast. And when the median falls, the poverty line falls with it — because the line is defined *as* a share of the median. A household earning exactly what it earned five years earlier could find itself reclassified from "poor" to "not poor," not because anything in its life had improved, but because the ruler measuring it had shrunk to match the collapse around it.

So AROP, the official income-poverty rate, barely moved through the worst years of the Greek debt crisis. Not because Greek households weren't getting poorer — they were, dramatically — but because the yardstick was falling with them, always re-centering on a poorer and poorer "normal."

What happens if you refuse to let the ruler move? Hold the poverty line to where it stood in 2008, adjusted only for inflation — not for how far the whole economy had fallen — and measured poverty roughly **doubles** during the crisis, from about 20% to a peak near 41%. That climb tracks almost exactly what Greek households were telling surveyors about their own lives, in real time, throughout those years.

Read the method and caveats

METHOD NOTE

Eurostat's own fixed-benchmark poverty product is anchored to 2019, not 2008, which misses the crisis years entirely. This report reconstructs an approximate 2008-anchored series: each year, Eurostat publishes the poverty rate at four income thresholds (40/50/60/70% of that year's median); fitting a probit-linear curve through those four points ($R^2 > 0.99$ every year) and evaluating it at the inflation-adjusted 2008 threshold recovers an estimate of "what share of people would fall below a 2008-level income."

The same method, applied to Eurostat's *real* 2019-anchored series and checked against the real published numbers for 2019–2025, is accurate to **1.2 points** on average (worst case 3.4 points) with a small, consistent upward bias and $r=0.89$ correlation to the true series. An independent academic reconstruction using real household microdata (Andriopoulou, Kanavitsa & Tsakloglou, 2019/2020) finds an even larger crisis-era effect than this approximation does — **48%** of the population poor by a 2007-anchor for income-year 2013, versus this report's **40.6%** for the aligned survey-year 2014. None of this is a new discovery. An OECD microsimulation study (Leventi & Matsaganis, 2016) finds the same arithmetic independently — Greece's official relative poverty rate barely moved across 2009–2014 even as absolute living standards collapsed, because the median it's measured against fell right along with everyone's income — and the OECD's own 2018 country survey makes the identical point. This chapter applies a well-established mechanism to Greece's subjective-poverty gap specifically; it doesn't claim to have found the mechanism itself.

A second Eurostat measure makes the same point from an entirely different angle, with no approximation involved at all. *AROPE* — "at risk of poverty or social exclusion" — counts anyone who is either income-poor, or severely deprived of basic goods, or living in a very-low-work household. Two of those three conditions have nothing to do with a shifting median. *AROPE* tracks Greek subjective poverty far more closely than *AROP* does, moving from 28% in 2008 to a peak of 36% in 2014 — a real, officially-measured rise, not a reconstruction.

Put the arithmetic fix and the official broader measure side by side, and a real, sizeable part of the Greek paradox dissolves. It was never entirely a mystery of national temperament. Part of it was always just a ruler that shrank at the worst possible moment.

What the ruler still doesn't catch

Fix the arithmetic, and a real gap survives anyway.

If the whole paradox were arithmetic, the story would end here: correct the yardstick, watch the mystery disappear. It doesn't. Compare Greece to other EU countries with similarly rough unemployment, income, and material hardship — a like-for-like comparison, not a Greece-only correction — and Greek households still report meaningfully more difficulty than those shared conditions alone would predict.

Add housing costs and cash-flow strain (missed bills, no buffer for a surprise expense) to that comparison, and most of the remaining gap closes — *if* the model is allowed to see Greece's own data while it learns. Test it more honestly — predict Greece using a model trained only on the other twenty-six countries, having never seen a single Greek data point — and the gap survives. Under every version of that stricter test built on headline unemployment — the standard comparison used throughout this section — Greece is the hardest country in the European Union to predict. (Chapter 7 finds a different labor-market measure that changes this picture substantially; this section describes what the standard comparison alone can and can't explain.)

Read the method and caveats

METHOD NOTE

In-sample gap (model sees Greece's own data): falls from 15.1 points (basic comparison: unemployment, income, deprivation, AROP) to 2.2 points once housing, arrears, and unexpected-expense capacity are added. AROP is used here, rather than AROPE, because AROPE partly bundles deprivation and work-intensity components that the model already measures separately. Out-of-sample gap (leave-Greece-out, stricter test): 25.6 → 35.5 → 11.6 points across the same three steps — housing alone makes the out-of-sample prediction *worse*, not better, before the fuller comparison brings it back down. Greece ranks 1st of 27 (largest gap) under each of these headline-unemployment specifications. One caveat worth carrying into the rest of this piece: two of those additions — being behind on

bills, and having no buffer for a surprise expense — are themselves close to the thing being explained. They measure financial strain, and they are being used to explain reported financial strain. That they close much of the gap is real, but partly definitional, which is why the chapters that follow lean on causes further upstream: how long people were out of work, and for how many years running.

That remaining 11.6-point gap is the real subject of everything that follows. Not a rounding error, not an artifact of a badly-chosen comparison year — a genuine, quantified residual that the standard toolkit of income, unemployment, and material hardship cannot explain. Something else is going on. The rest of this piece is that something else, told as what it actually is: not one hidden variable, but a sequence of scars that never fully closed.


The economy that didn't come back

Most countries that fell, got back up. Greece was hit twice before it had healed from the first fall.

Almost every EU country that suffered a genuine crisis-era downturn recovered from it within a few years — a median of three, and none took longer than seven except Greece. Greece's own path had two separate chapters, not one uninterrupted slide. The debt crisis alone drove a 26.2% decline from the 2008 peak to a first trough in **2013**; real recovery followed, closing to within 20 points of the old peak by 2019. Then the pandemic cut that recovery off, pushing the index to a new low — 26.9% below peak — in **2020**, barely worse than the original 2013 trough, but from an unrelated shock landing on an economy that hadn't yet finished healing from the first one, not from twelve straight years of falling. Recovery resumed again afterward. As of the most recent data, Greek real GDP per capita remains 11.2% below its 2008 peak, sixteen years on.

Two smaller countries complicate the "only" in that sentence, and it is worth naming them rather than smoothing them away: Finland came within 0.8% of its own 2008 peak in 2022, then slipped back below it; Luxembourg currently sits 6.1% below a peak of its own. Neither changes the headline. Luxembourg's dip is small, recent, and still unfolding in an unusually wealthy economy — a different kind of story from a sixteen-year-old unhealed crisis. By scale and by duration, Greece's decline is still, by any reasonable measure, unresolved in a way no other EU economy's is.

COUNTRY	% BELOW ITS OWN 2008–2024 PEAK, 2024
Greece	11.2%
Estonia	7.1%
Luxembourg	6.1%
Ireland	3.3%
Austria	2.8%
All other EU countries (18 of 27)	0.0–2.5%, mostly fully recovered



This is where the report's real argument begins to take shape. Not "Greeks compare themselves to an abstract European average and feel bad about it" — something more specific and more falsifiable: the economy itself never gave most households their old life back. Everything in the chapters that follow is a different, more concrete way of watching that same failure to heal.


A smaller paycheck

GDP is an abstraction. The number on a payslip isn't.

Real GDP per capita is an economist's number — a useful one, but not something any household actually receives in an envelope. The paycheck is. And the paycheck tells the same story a second time, more directly.

–31.8%

How far below their 2008 level Greek **real wages** still sat in 2024 — the largest shortfall of any EU country, and not close. Hungary is a distant second at –25.8%.



Real wages (compensation per employee), indexed to each country's own 2008 level, 2008–2024. Capped at 160 to keep the Greece-comparison zone readable — a few Eastern European countries (visible flattening at the top) grew well past that line. Source: Eurostat.

Real wages — compensation per employee, adjusted for inflation — peaked in Greece in 2009 and bottomed out in 2023. By 2024 they remained 31.8% below 2008, the EU's largest shortfall; Hungary was second and Italy third, while most countries had recovered or grown well beyond their old level. Greece is not the only country still below 2008. It is the country with the deepest unresolved paycheck loss.

Read the method and caveats

METHOD NOTE

Built from nominal compensation per employee (Eurostat `nama_10_lp_ulc`), deflated by the national HICP price index, rebased so each country's own 2008 value equals 100. At the level, real wages track subjective poverty about as strongly as any variable in this report ($r=-0.79$) — but on year-to-year changes that relationship is weak and not statistically significant, the same pattern this report flags for several level-only variables. Tested directly as a candidate addition to the main cross-country model: it does not add independent explanatory power once existing predictors are accounted for, so it is kept here as strong descriptive evidence, not folded into the model.

Ordinary things, extraordinary weight

Greece is not an expensive country. That is exactly the point.

Here is a number that seems, at first, to contradict everything so far: on raw price level, Greece is not one of the EU's expensive countries. Its overall cost of living ranks only 18th-highest of 27 — below the EU average, not above it. Denmark is expensive. Ireland is expensive. Greece, by the numbers a tourist would recognize, is not.

But a price level only tells you what things cost. It says nothing about what a household can afford to spend. And once you divide those Greek prices by a Greek paycheck — comparing what a euro of Greek wages actually buys, against what a euro of the EU-average wage buys — the ranking does not just shift. It **inverts completely**.

18th → 1st

Greece's rank among 27 EU countries on overall consumption: **18th-most-expensive** by raw price level, but **1st** for wage-adjusted price pressure — the highest in the EU.

The reason is simple, and stark: Greek wages sit at just 48.3% of the EU average — the 4th-lowest in the Union. Prices don't need to be high for the squeeze to be severe when the paycheck behind them is this small. The pattern holds category by category, not just in the aggregate — and it holds hardest in exactly the categories where Greek prices look, on paper, unremarkable.

CATEGORY	RAW PRICE (EU=100)	WAGE-ADJUSTED PRESSURE	EU RANK, PRESSURE
Overall consumption	83.6	173.1	1st of 27
Food & beverages	106.2	219.9	2nd of 27
Housing, water, electricity, gas	73.5	152.2	2nd of 27
Transport	89.7	185.7	1st of 27
Restaurants & accommodation	86.6	179.3	1st of 27

And the two directions of that squeeze — the shrinking paycheck, the rising bill — have not stood still since 2008. They have moved apart. Greek nominal wages, unadjusted for inflation, sit 13.1% below their 2008 level even today. Over that same span, nominal food prices rose 41.6%. Housing and energy rose 33.6%. Greek households do not need to imagine hardship to feel squeezed. The receipt grew while the paycheck shrank, in absolute euros, at the same time.

Read the method and caveats

METHOD NOTE — A UNITS CLARIFICATION

Two different comparisons appear across this section and the previous one, and they measure different things: real wages (Chapter 4) are *inflation-adjusted purchasing power*, tracked against Greece's own 2008 level over time. Wage-adjusted price pressure here is a different, *nominal* comparison — actual euro wages against actual euro prices, benchmarked against the rest of the EU rather than against Greece's own past. Both are correct; they answer different questions. This is a comparative price-level index (Eurostat's harmonized category comparisons), not an actual household-budget survey or real receipts — and the category-level detail has full 27-country coverage for only 2022–2024, too short a window to support a formal cross-country model test. It is reported here as strong descriptive evidence, deliberately labeled "price pressure," not "affordability," since it is a transparent ratio of two published indices rather than a constructed cost-of-living methodology of its own.

No longer safe even owning it

In most of Europe, paying off your mortgage means the housing squeeze is over. Not here.

In most EU countries, owning a home outright is close to a guarantee: pay off the mortgage, and housing largely disappears as a source of financial strain. That protection is one of the more dependable patterns in European household economics — which is exactly what makes Greece's version of it so telling.

25.7%

Share of Greek **mortgage-free homeowners** still overburdened by housing costs in 2024 — the highest in the EU, and almost **twice** the next-highest country, Sweden (14.0%). The typical EU country: 1–7%.

Renters are still worse off in absolute terms — 37.4% of Greek market-rate renters are overburdened, against 25.7% of mortgage-free owners. But look at the *size of that gap*, not just the two numbers, and something more unusual appears: Greece's owner-renter divide is one of the narrowest in the EU, ranking 22nd of 27 on the size of the difference. Not because renters are spared — they aren't — but because ownership, which shields households almost everywhere else in Europe, barely shields them here. Greece's housing problem does not stop at the renter-owner line the way it does elsewhere. It reaches people who, on paper, should be protected.

The ownership structure itself carries a clue about why. Greece is not unusual in how many people own their homes — roughly 70%, close to the EU average. It is unusual in *how* they own them: only 8–17% of Greek ownership is mortgage-financed, against a stable 24–26% EU-wide. Most Greek homeowners hold their

houses outright, likely reflecting an older, inherited, or self-built housing stock rather than credit-financed purchase — consistent with, though this report does not directly test, pressure from property tax, aging infrastructure, and energy costs rather than rent or mortgage payments themselves. One further, still-unfolding shift is worth flagging: Greece's ownership rate has fallen faster than the EU average recently, from 73.9% to 69.7% between 2020 and 2024 — more than double the EU's own pace of decline over the same years.

Read the method and caveats

METHOD NOTE

Source: Eurostat housing cost overburden by tenure status (`ilc_lvho07c`) and tenure distribution (`ilc_lvho02`), full 27-country coverage. The "TOTAL" tenure row exactly matches the general housing-cost-overburden variable already used throughout this report's cross-country model — this chapter is a drill-down into a variable already central to the story, not a new independent one. Mortgage-free-owner overburden correlates at $r \approx 0.9$ with that existing variable (expected, since it is a subcomponent of the total), and adding it as a further model predictor changes little and is not statistically significant — kept as descriptive depth on the housing-cost mechanism, not a new model finding.


When joblessness stops being a spell

The single strongest signal in this entire investigation was never about how many people were unemployed. It was about how long.

Every measure so far has described conditions that stayed bad. This one describes the mechanism that best explains *why* they stayed bad — and it is the one result in this entire report that actually changes the model's verdict, not just its texture.

Headline unemployment tells you how many people are out of work at a given moment. It says nothing about whether the same people have been out of work for a month or for years — whether the labor market is actually re-absorbing them, or simply cycling them through. *Long-term unemployment* — the share of the labor force out of work for twelve months or more — answers that different question. And it turns out to be the strongest single variable in this entire investigation, at every time scale tested: how high it runs, how much it moves year to year, and how far it deviates from its own trend all track Greek subjective poverty more tightly than anything else measured.

Greece's long-term unemployment rate has fallen a long way from its worst point — from 17.5% of the labor force in 2014 (meaning roughly three-quarters of all unemployed Greeks, at that moment, had been out of work a year or more) down to 5.4% in 2024. That is real, substantial progress. It is also still, by a wide margin, the highest rate in the European Union: next is Spain at 3.8%, then Slovakia and Italy near 3.5%. Most of the EU sits under 2%. Headline unemployment recovered far faster than this did.



Long-term unemployment rate (12 months or more, % of active population), all EU countries, 2009–2024. Source: Eurostat.

Swap this one measure into the model in place of headline unemployment — changing nothing else — and Greece's story changes with it.

PROOF TABLE

What changes when the model uses duration instead of rate

Same predictors, same panel, same method — only the unemployment measure changes

MODEL	FIT (R ²)	GREECE'S GAP, OUT-OF-SAMPLE	GREECE'S RANK
Headline unemployment (baseline)	0.89	11.6	1st of 27
Long-term unemployment (replacing headline)	0.91	3.9	6th of 27

Without adding a single new variable to the model, this one substitution cuts Greece's unexplained gap by roughly two-thirds and removes it from the top of the outlier list entirely. Full 27-country ranking, before and after — Cyprus, Luxembourg, Portugal, Czechia, and Belgium now rank above Greece.

That does not make the underlying hardship disappear. It sharpens what caused it. The crisis did not simply put people out of work for a while. It left a meaningful share of them out of work long enough for the experience to stop being a bad year and start being a condition — and a model that can see that distinction understands Greece far better than one that can't.

Read the method and caveats

METHOD NOTE

Long-term and headline unemployment are, unsurprisingly, highly correlated ($r=0.91$ panel-wide; $r=0.998$ for Greece's own series, since by the mid-2010s most of Greece's unemployment was long-term). Adding both to the same model together destabilizes headline unemployment's own coefficient — a textbook collinearity symptom — so replacement, not addition, is the preferred specification. Long-term unemployment's own coefficient is stable and significant ($p<0.001$) across all 27 leave-one-country-out refits, including the refit that excludes Greece itself: not an artifact of Greece's own data. Being steady is not the same as being precisely measured, though. Under the stricter test appropriate to 27 countries, long-term unemployment is measured loosely; the accumulated total of unemployment is not. The firmer of the two findings is the accumulated one. Two findings from earlier chapters turn out to overlap with this one once tested together — the GDP-scarring result adds nothing further once long-term unemployment is in the model, and financial pessimism weakens from significant to marginal — addressed honestly in Chapter 11.

The debt that never gets paid down

Long-term unemployment says how many people are stuck right now. It doesn't say how long the whole country has been carrying this.

A rate is a snapshot. It answers "how bad is it this year," and nothing about how many years came before it. Two countries can arrive at the same unemployment rate today by very different roads — one that had a brutal year and mostly recovered, and one that has had a decade of mildly-bad years and never recovered at all — and a snapshot cannot tell them apart. This chapter asks a cumulative question instead: not "how bad is it now," but "how much has this economy made people carry since the crisis began, added up, year after year, and never subtracted back."

The construction is simple and deliberately unforgiving. For every country, starting in 2009, take that year's unemployment rate minus the country's own 2009 baseline. If the country is doing better than its own baseline, that year contributes nothing — not a negative number, not a credit toward a future bad year. Just zero. If it's doing worse, the excess gets added to a running total that only ever grows. A good year can't cancel a bad one that already happened; it can only stop the bleeding from that point forward. The result, for every country, is a single number: how many "excess unemployment-years" it has accumulated since the crisis began.

+2.7pt

Greece's residual under the fully nested select-then-predict procedure: 19th of 27 on prediction error, with eight countries predicted less accurately. The fixed cumulative model reaches -0.8 ; $+2.7$ is the more conservative result because it includes the cost of choosing the variable.

This is the single strongest addition tested anywhere in this investigation. Not because it's a clever new variable pulled from nowhere — it's built entirely from unemployment data already in the model — but because of what it captures that a snapshot cannot: that Greece didn't just have a bad labor market at various points since 2009, it has been continuously carrying the accumulated weight of one, year on top of year, for longer than almost anywhere else in Europe.

PROOF TABLE

The gap-closing ladder, from the raw puzzle to the final model

All rows share one 2015–2024 window. Steps 1–2 are raw, single-country gaps; steps 3–6 are out-of-sample model residuals — two different kinds of quantity, a sequence rather than one additive decomposition

STEP	WHAT'S ADDED	GREECE'S GAP
1. Raw AROP gap (2015–2024 avg)	—	52.6
2. Raw AROPE gap (2015–2024 avg, the bridge)	broader official measure	41.5
3. Basic model (Chapter 2)	unemployment, income, deprivation	25.6
4. + housing, arrears, expenses (Chapter 2)	cash-flow strain	11.6
5. Headline → long-term unemployment (Chapter 7)	duration, not just rate	3.9
6. + cumulative excess unemployment (this chapter)	accumulated exposure since 2009	–0.8

The table shows the fixed-model sequence. Full nested validation then repeats selection inside every fold and places Greece 19th of 27 on prediction error. Both results remove Greece from the exceptional-outlier position; the nested result is the stronger test.

The result holds up under scrutiny, not just at face value. Refit the model 27 times, excluding one country at a time, and cumulative excess unemployment's own coefficient never changes sign and never loses significance — including the refit that excludes Greece itself. That's a check on whether the *coefficient* is stable once this is the variable being tested, though — a separate question, addressed below, is whether Greece's own data should have counted toward *picking* this variable over its closest rival in the first place. It is not a substitute for long-term unemployment, either: replacing long-term unemployment with this cumulative measure, rather than adding it, performs worse. The two are doing different jobs — one measures who is stuck right now, the other measures how long the whole country has been carrying the weight — and both matter.

METHOD NOTE

Baseline year 2009 (the earliest year with near-complete 27-country coverage for both headline and long-term unemployment, confirmed by direct check). Excess above baseline, floored at zero, summed year over year through 2024. Added to the Chapter 7 model: R^2 rises from 0.914 to 0.930; Greece's out-of-sample gap moves from 3.9 to -0.8. Leave-one-country-out coefficient range: 0.134–0.197, always positive, always significant, including excluding Greece (0.164, $p=0.0064$). Replacing long-term unemployment with this measure instead of adding it performs worse (R^2 0.911, gap 4.3) — confirming it's a genuine additional layer, not a relabeling. A related measure — consecutive years Greek real wages stayed below their own 2008 level — independently supports the same story ($p=0.0045$, stable under the same leave-one-out check), though within Greece's own data the two move too closely together (correlation 0.95) to credibly count as two separate mechanisms; kept as supporting evidence. One anticipated candidate was tested and did *not* hold up: a cumulative shortfall in the AROP poverty threshold itself, despite being the most obvious extension of Chapter 1's shrinking-ruler story, came back statistically null ($p=0.291$).

A negative out-of-sample gap means the model's prediction — refit on the other 26 countries, with this specific variable already chosen and Greece's own data point held out of that refit — runs slightly ahead of what Greek households actually report, not that Greek households are secretly optimistic. It is a property of this specific, richest model, not a claim that the whole puzzle is now "more than 100% explained," a phrase this report deliberately avoids. Checked year by year, the fit is close for nine of the last ten years (residuals from -3.3 to +1.0); the one exception, 2021, plausibly reflects a pandemic-specific disruption the model isn't built to capture — not a sign the overall fit is a fluke.

On multiple testing: eighteen cumulative and duration candidates were screened as one family in this chapter (the original 8 plus the 10-variable duration battery), and the correction is computed inside the same script that ran the screening, not as a separate afterthought. Checked against Benjamini-Hochberg false-discovery-rate correction, cumulative excess unemployment survives comfortably (adjusted $p=0.0024$); the wage-duration measure survives, but narrowly (adjusted $p=0.0408$, just inside the usual 0.05 cutoff); none of the other 16 candidates does. This confirms this chapter's treatment of cumulative excess unemployment as the headline result and the wage-duration finding as supporting evidence held to a visibly lower bar — a call already made on magnitude and stability grounds before this correction was run, not decided by it. What that correction does *not* do is make this a pre-registered test. These models were built up over a sequence of checkpoints, each informed by the last, so the honest description is a well-validated result arrived at by searching — not a hypothesis stated in advance and then confirmed. That is exactly why the weight here rests on the out-of-sample and leave-one-country-out checks rather than on any single p-value.

An honest complication: which of the two was really "picked" depends on whether Greece counts. Cumulative excess unemployment was selected as the headline variable using the full 27-country panel, Greece included — a legitimate way to screen candidates,

since Greece's data is a real part of the historical record, but worth being explicit about. Rerunning that same 18-candidate screening with Greece dropped from the panel entirely, not just held out for evaluation, the wage-duration measure edges ahead ($p=0.0036$ vs. $p=0.0064$, both still highly significant). The two candidates that matter don't change; which one ranks first does. Best read as two closely related, mutually reinforcing signs of sustained hardship, not one uniquely identified mechanism with a clear runner-up. (A stricter version of this check was run later: the whole 18-candidate screening was re-run once for every possible excluded country, not just Greece. The same variable wins in 25 of 27 re-runs, and both exceptions pick closely related duration measures. Stricter still, each re-run's own winner was then used to predict the country that had been left out of it entirely — and on that fully arm's-length test Greece comes out mid-pack, 19th of 27, with eight EU countries predicted *less* accurately than Greece. Greece's gap under that test is $+2.7$ points rather than the -0.8 of the fixed final model — a more conservative number, because it also pays the cost of choosing the variable in the first place. Full detail in the technical report's Methods.)

Is this really "permanent," or just "sustained"? Because the running total can only grow, on its own it can't tell "genuine undiminished scarring" apart from simply counting years since the crisis started. Tested five alternatives that *can* fall over time — trailing 3-, 5-, and 10-year rolling sums, and two versions that let old years fade at 20%/year and 10%/year. A 3-year window misses the effect entirely; a 10-year window fits at least as well as the permanent since-2009 sum, slightly better by some measures. So the honest version of this chapter's claim is sustained excess unemployment over roughly the past decade, not literal permanent accumulation since one specific year — a claim now backed by several related constructions pointing the same way, not resting on one arbitrary choice.

Not without precedent, and not without limits. Eurostat already applies a version of this logic at the individual level: its own persistent-poverty measure counts someone poor only if they were poor this year *and* in at least two of the past three, precisely because a single-year snapshot is known to miss how widespread and lasting poverty experience really is — a 2026 European Commission study of EU-wide longitudinal data reaches the same conclusion and calls for more dynamic measures in policy design. Cumulative excess unemployment applies the same idea one level up: a country's labor market, not a single household, carrying accumulated exposure rather than a snapshot. At the individual level, research following unemployed people directly over many years finds something similar — life satisfaction doesn't fully bounce back even after re-employment, and each extra spell of past unemployment predicts lower wellbeing years later, even accounting for current conditions. But those studies track individual people; this chapter's variable tracks a whole country's economy. The correlation between Greece's accumulated unemployment exposure and its subjective-poverty rate is consistent with the same underlying mechanism, but it isn't direct proof that the specific households behind that 67.2% number personally lived through years of joblessness. That gap between country-level and household-level evidence is a real limit of what this chapter's data can show.


Some people just left

Income and jobs are not the only things a household can lose to an unrecovered crisis. People themselves can leave.

This is not a hidden variable in any model — it is the visible, human consequence of everything in the eight chapters before it. During exactly the years income, jobs, and expectations were collapsing, a meaningful share of Greek households treated leaving the country as the answer.

290,281

Net loss of Greek nationals, 2008–2024 — 2.6% of Greece's 2008 population. Outflow peaked at **44,502** in a single year, 2012, the most acute stretch of the crisis.



Net migration of Greek nationals by year. Rust bars: more people left than returned that year. Green bars: more returned than left. Source: Eurostat.

The outflow has, only very recently, begun to reverse. 2023 and 2024 both show a net *inflow* — more Greek nationals returning than leaving, for the first time since before the crisis began. That reversal is real. It does not erase the cumulative loss, or the years of disrupted careers and family separation behind it, and Greece is not even the EU's most extreme case on this specific measure — several Eastern and Baltic countries lost a larger share of their population to emigration over the same period. What makes Greece's version notable isn't uniqueness. It's that the outflow was large, crisis-linked, and sustained for well over a decade, arriving at exactly the moments the rest of this story describes.

Read the method and caveats

METHOD NOTE

Source: Eurostat migration-flow tables (`migr_em1ctz` / `migr_imm1ctz`), Greek nationals only, 2008–2024. Youth unemployment (ages 15–24) peaked at 59.2% in 2013, one year after the emigration peak — a plausible, but not causally tested, push-factor timing. Youth unemployment tracks subjective poverty strongly on its own, but overlaps so heavily with long-term unemployment that it adds no independent explanatory power once that variable is already in the model — kept here as descriptive context, not its own finding.

The national average hid a generational reversal

Every number so far treats Greece as one population. It isn't. Break the same official measure down by age, and "the aggregate got a little worse in 2025" turns out to be the wrong headline entirely.

AROPE, AROP, and material deprivation are all published broken down by age group, not just as one national figure. This isn't a new candidate explanation for the cross-country gap the rest of this piece traces — it doesn't enter any model here — it's a distributional finding about who inside Greece is actually carrying the number, sitting alongside the AROPE bridge from Chapters 1–2, not on top of the mechanism in Chapters 7–8.

+0.65pt

People 65+ alone contributed more to Greece's 2024→2025 national AROPE increase than the entire net change — while every working-age group's own rate *improved*.

PROOF TABLE

AROPE by age group, Greece: 2015 versus 2025

Every working-age group moved down. People 65+ moved the other way.

AGE GROUP	2015	2025	CHANGE
Under 18	37.7%	29.6%	−8.1pt
18–24	43.6%	33.2%	−10.4pt
25–49	35.0%	25.1%	−9.9pt
50–64	34.8%	27.5%	−7.3pt
65+	17.5%	27.7%	+10.2pt

Young adults remain the single highest-risk group in absolute terms (33.2%) — but that's a recovery from an exceptionally bad starting point, not a decline. People 65+ started this decade as the *lowest*-risk age group in Greece and are now higher than every working-age group except the youngest.

Why did the national number rise at all, if four out of five age groups improved? Not because Greece's population is aging into a higher-risk bracket — that effect, measured directly, is almost exactly zero. It's because the 65+ group's own rate rose so much that it outweighs everyone else's improvement combined.

PROOF TABLE

Why the national rate rose in 2025: a shift-share decomposition

Splitting the +0.6-point national change into "the population shifted toward older groups" versus "each group's own rate changed"

EFFECT	CONTRIBUTION
Population aging (composition shift)	+0.02pt
Each group's own rate changing	+0.60pt

Broken down further, by group's own contribution to that +0.60pt: under 18 +0.28pt, 18–24 –0.09pt, 25–49 –0.09pt, 50–64 –0.15pt, 65+ **+0.65pt**. The elderly contribution alone exceeds the entire net national change.

Among people 65+, the deterioration isn't spread evenly either. It concentrates specifically on women and people living alone.

PROOF TABLE

Who among people 65+ is most exposed

GROUP	2015	2025	CHANGE
Women 65+	18.7%	30.9%	+12.2pt
Men 65+	15.9%	23.6%	+7.7pt
65+ couples (2 adults, at least one 65+)	18.2%	26.2%	+8.0pt
Single-person households, 65+	23.0%	38.7%	+15.7pt

At EU level over the same decade, both groups rose only modestly by comparison (women 20.6% to 21.2%, men 14.7% to 15.8%) — Greece's elderly women are diverging sharply from the EU pattern, not simply worse off in the same direction. Greek single-elderly households started 2015 *below* the EU average (23.0% vs. 26.9%) and are now *above* it (38.7% vs. 30.0%): a genuine reversal, not a bad starting position falling further behind in step with everyone else.

METHOD NOTE

Why this chapter's AROP numbers could go back to 2003, but its AROPE and deprivation numbers don't. Age-specific AROP is a methodologically continuous series — no legacy/revised split exists for it, so it can be shown from 2003, same as the rest of this report. Age-specific AROPE and material deprivation are shown from 2015 only, because Eurostat's revised age-breakdown series begins there, and older legacy values aren't spliced in: a direct check of the overlap years (2015–2020, where both legacy and revised data exist) found the age-specific gap between the two definitions too large to trust — up to 7.3 points for 18–24, 5.4 points for 65+, well past the roughly 1–3-point gap this report already tolerates when splicing the whole-population AROPE series elsewhere. For context, not as this chapter's main point: Greek 65+ AROP was 29.4% back in 2003, fell steadily to a low of 11.6% in 2018, then climbed for seven straight years to 20.9% in 2025 — still below where it started in 2003. The current elderly reversal is real, but it's a partial undoing of a larger earlier improvement, not new, unprecedented territory. That longer arc is background; the 2015–2025 window above is this chapter's real comparison, since it's the only span where AROP, AROPE, and deprivation can all be read together on one consistent basis.

Age groups follow Eurostat's own standard EU-SILC breakdown, not a custom bucketing. Household figures use Eurostat's own "one adult 65+" and "two adults, at least one 65+" categories directly. Three things were checked and deliberately left out, rather than approximated: sex within single-elderly households (Eurostat's household tables have no sex dimension at all, and the single-male/single-female categories elsewhere aren't age-restricted, so combining them would fabricate a cross-tab that doesn't exist); tenure by age (the housing-overburden-by-tenure table has no age breakdown, and the one table that does cross household type with tenure measures population *distribution*, not an overburden *rate* — not the same thing); and a formal margin of error (Eurostat's API exposes none for these cells, checked directly rather than assumed). The practical stand-in: Greece's single-65+ population is real but modest — 273,000 people as of 2025 — so single-year movements for this and smaller subgroups carry more normal survey variance than the national headline. The decade-long trend is the load-bearing claim here; the sharp 2024→2025 jump corroborates it, not proves it on its own. And very-low-work-intensity, AROPE's third component, has no 65+ breakdown at all in Eurostat's data — structurally undefined for this age group — so whatever is driving the 65+ increase has to be income poverty, material deprivation, or their overlap, not work attachment.

Not hardship restated. Hardship extrapolated.

Pessimism here isn't a household reporting that things are hard. It's a household that has watched things stay hard for sixteen years, concluding there's no particular reason to expect that to change.

Greece is the single most pessimistic country in the European Union about where its own households' finances are headed — worst of 35 countries surveyed, 11 points clear of the next-most-pessimistic, Slovenia. On its own, added to the baseline model, that pessimism measure closes real ground: it knocks Greece from the largest cross-country gap in the EU down to 10th of 27.

The easy, and wrong, reading is that this is circular — households asked about tomorrow are just restating how bad today already is, so using one to explain the other explains a feeling with itself. There's a real version of that concern: once long-term unemployment (Chapter 7) is in the model instead of headline unemployment, financial pessimism's own contribution weakens from statistically significant to marginal, and some of what looked like an independent psychological effect turns out to overlap with the labor-market mechanism already established. But the more precise reading isn't "restating the present." It's *extrapolating from the past*. A household seventeen years into an economy that has never fully come back does not need to imagine hardship continuing — it has already watched hardship not end, past every point at which a comparable European economy's did. Expecting no improvement, under those specific conditions, isn't catastrophizing. It's closer to pattern recognition.

The political backdrop gives that pattern nothing to push back against. Greek institutional trust was never particularly high before the crisis (averaging 44% through the 2000s), collapsed to a historic low of 7% in 2012 — the same year this

report's own migration data marks as the acute peak of the crisis — and has recovered only partially since, to 37% at best. It is falling again as of the most recent data, despite genuinely strong recent macroeconomic performance: the one piece of unambiguous good news in this entire investigation, and it has produced no corresponding rebound in how much people trust the institutions meant to manage the economy on their behalf. The trust series isn't tested as a model mechanism anywhere in this report — it's context, not a variable that explains anything away — but it points at the same shape as the income story: a long stretch where the thing that was supposed to recover didn't, on schedule or otherwise. That parallel helps explain why a macro recovery can arrive in the statistics without yet arriving in how a household feels about the years ahead.

Read the method and caveats

METHOD NOTE

Financial expectations (Eurostat household consumer survey): added to the headline-unemployment model, R^2 rises $0.892 \rightarrow 0.904$, coefficient significant ($p=0.021$). Added to the long-term-unemployment model instead: coefficient weakens to $p=0.062$, no longer significant at the conventional threshold. Household saving rate was also tested as a wealth-depletion signal (Greece 2nd-lowest of 34 countries) but added no independent power once expectations were already in the model. Institutional trust itself is not modeled — Eurostat's own holdings on the topic amount to a single year of data (2013), too thin to support a trend variable — and is included here only as literature-backed interpretive context (Ervasti, Kouvo & Venetoklis 2019; OECD 2026 trust survey). The trust-and-income-recovery timelines are presented here as a documented parallel, not a tested joint mechanism — no model in this report or the technical report formally links them.

Working the most, earning the least

Everything so far has treated unemployment as the fault line. Here's the sharper test: take joblessness off the table entirely. Does having a job actually fix this?

For Greek salaried workers specifically, no. On the official income-poverty measure, salaried Greeks look almost boringly average — 5.7% AROP, ranking 17th of 27 EU countries, squarely in the middle of the pack. And yet 59.5% of them — nearly six in ten — say they struggle to make ends meet, against an EU average of 13.6%. Still first in Europe, by a wide margin, on the one measure that asks people how it actually feels. This is the same paradox this whole piece opened with, but with unemployment subtracted out completely. It isn't a story about people without jobs. It's there among people with them.

One obvious next question: maybe it's not about how much Greeks earn, but how much they have to work for it. It turns out to be both, at once. Greece works the longest hours of any EU country — 39.8 a week across all employed people, highest in the Union, and still highest restricted to full-time workers alone (41.1). Restricted to salaried employees only, Greek hours are high rather than extreme — 7th-highest of 27 — so the self-employed contribute to the record, but they aren't the whole story. And for each of those hours, Greek workers are paid the least of anyone in the EU: €14.2-equivalent per hour in purchasing-power terms, against an EU average of €29.6. Put the two together — hours worked relative to the EU, divided by pay per hour relative to the EU — and Greece's score is more than double the EU average, the highest of any member state. It isn't a coincidence that the country that works the most also earns the least per hour of it. It's the same imbalance, seen from two directions.

The pay side of that imbalance has also lost ground over time. Real, inflation-adjusted compensation per hour worked in Greece sits 27.8% below its 2008 level. So the country most reliant on long hours to make ends meet is also the country

where an hour of work now buys noticeably less than it used to. Working harder has not been enough to keep pace with what work itself used to be worth.

One honest caveat, checked rather than assumed: this is a description of who Greece *is*, not a story about what changed for Greece from one year to the next. Tested the way the strongest findings elsewhere in this piece were tested — does the measure track a country's own year-to-year movement, not just its position relative to its neighbors — it fails decisively. Add country-level controls and the effect disappears (not statistically significant). Look at year-over-year changes instead of levels, same result. And within Greece's own fifteen-year history (2010–2024), this squeeze measure and subjective poverty are essentially uncorrelated. Long-term unemployment and its accumulated version (Chapters 7–8) explain *movement* — they track Greece's own worsening and improving, year by year. This finding explains something different: a stable, structural fact about where Greece sits relative to the rest of Europe, present throughout, not something that got worse and then better on its own timeline. Both kinds of evidence matter. They're just not the same kind.

Read the method and caveats

METHOD NOTE

Salaried-worker AROP, AROPE, and subjective poverty: Eurostat's in-work poverty and subjective-poverty-by-status tables, 2025. Hours: LFS actual weekly hours in the main job (`lfsa_ewhan2`), 2024, the latest year with matching hourly-compensation data. Hourly compensation: national-accounts compensation per hour worked, purchasing-power-standard adjusted (`nama_10_lp_ulc`). Added to the strongest existing cross-country model, the work-effort squeeze survives correction for testing multiple candidates together (FDR-adjusted $p=0.003$) and moves Greece's out-of-sample gap from +3.9 points to -3.5 — but country fixed-effects ($p=0.34$) and year-over-year first-differencing ($p=0.73$) both fail, and Greece's own time-series correlation is $r=-0.03$. Kept here as supporting structural evidence, not added alongside long-term or cumulative unemployment as a dynamic mechanism. Full derivation: `scripts/43_work_effort_squeeze.py`.

What wasn't it

A useful investigation also reports what it ruled out.

Not every plausible suspect held up. Reporting the negatives honestly is part of what makes the positives worth trusting.

RULED OUT

Income inequality. Greece's inequality is elevated but not exceptional — 6th-highest of 27 EU countries on both the Gini coefficient and the S80/S20 ratio — and its own inequality has actually *fallen* since 2003, even as subjective poverty stayed elevated. Added to the model, it explains none of the remaining gap. This report's poverty paradox is not a story about a hidden unequal tail; it is a story about average-level hardship.

TESTED, NULL

A cumulative shortfall in the AROP poverty threshold itself. The most narratively obvious extension of Chapter 1's shrinking-ruler story — tracking how far Greece's own poverty line fell behind its pre-crisis path, accumulated over time — and it came back statistically null (Chapter 8). A reasonable hypothesis the data simply didn't support.

REDUNDANT

Youth unemployment. Genuinely strong on its own, but so heavily overlapped with long-term unemployment that it adds nothing once that variable is already in the model. Kept as context (Chapter 9), not counted twice.

REDUNDANT

The GDP-scarring stock. A real, standalone finding (Chapter 3) — but once long-term unemployment is in the model, its further contribution is no longer statistically significant. Long-term unemployment appears to capture the same underlying scarring more directly.

TESTED, NOT SCORECARD

Household debt and Greece's welfare-transfer system.

Tested directly; provided no additional explanatory power in this specification — a narrower finding than ruling them out, since their effects could be tangled up with variables already included.

Is this just how Greeks talk?

*A fair question, and not a comfortable one:
what if the whole premise of this piece is a
matter of tone, not fact?*

Subjective poverty is, after all, a self-report. It asks a household how hard it is to make ends meet, not what it earns or owns. If Greek respondents simply describe hardship more darkly — a harsher cultural baseline, a bleaker way of answering a survey — then some of what this piece has spent thirteen chapters explaining could be a matter of dialect, not distress. That possibility deserves a direct answer, not a wave of the hand, so here is one: six separate tests, run rather than assumed, none of them softened after the fact to fit a preferred conclusion.

Start with the uncomfortable part, conceded rather than argued away: Greece really was already among the EU's highest on this measure before the crisis had even begun. Averaged over 2003 to 2008, its subjective poverty was near the top of Europe. That's a real finding, and worth being honest about — except the years it's set against turn out to matter more than they first appear. The EU-wide survey behind this whole piece launched in 2003 with only six countries fully on board: Belgium, Denmark, Greece, Ireland, Luxembourg, Austria — a detail confirmed directly from the survey's own documentation, which separately notes real disruption in the data as later countries joined through 2005. Compare Greece against a proper, apples-to-apples panel of those original six, rather than an incomplete stand-in for the whole EU, and the ranking doesn't soften. It sharpens: Greece was 1st, not 2nd, among the countries the survey actually covered from day one. A pre-existing baseline is real. The question is whether it's the whole story.

It isn't. A stable cultural trait should produce a stable gap — roughly as wide before the crisis as after. Greece's gap between what households report and what the official income-poverty rate shows was 30.4 points in 2003–2008. By 2019–2024 it was 50.2 — a widening of nearly 20 points, the largest of any EU country, while the typical other EU country's gap over the same span actually *narrowed*.

Tested formally — Greece's own shift after 2009 against every other country's shift over the same years, with each country's baseline and each year's EU-wide shock already stripped out — that widening survives everything thrown at it: it doesn't depend on exactly which year the crisis is dated to; it barely moves if any single comparison country is dropped; it's if anything larger against Greece's closest peers alone (Italy, Spain, Portugal, Cyprus, Malta); there's no sign the gap was already drifting wider before 2009, which would have suggested Greece was simply on a different path all along; and pretending, in turn, that each of the other 26 EU countries was "Greece" and re-running the same comparison, Greece's own shift is still the largest of all 27 — one in twenty-seven, or 0.037. That number is worth reading carefully. It is not the output of an experiment: nobody assigned the crisis to Greece and withheld it from the others. It says that if the other twenty-six countries are treated as fair stand-ins for Greece, only one arrangement out of twenty-seven looks as extreme as the real one. The count is exact; treating it as a probability is the part that requires the assumption. One further check was tried and did not work, which is worth saying plainly. The idea was to build a stand-in Greece — a blend of other countries chosen to match Greece before the crisis — and then watch the real country and its stand-in come apart. Over the years that would make this convincing, 2003 to 2008, only six countries have complete coverage throughout — others report in some years but not others — and no blend of them looks like Greece: the stand-in sits around three to eight points where Greece sits above thirty. Something that never resembled Greece to begin with cannot tell you much by drifting away from it later. Narrow the match to just two years and it fits almost perfectly, but two points and twenty-five dials can readily be made to line up, so that agreement counts for very little. So this check is set aside. What stands is the before-and-after comparison and the placebo test.

The sharpest test is the simplest one. If this were really just a Greek way of talking, it should show up everywhere Greek households are asked to describe their own lives — not only on questions about money. It doesn't. On "can you make ends meet," Greece ranks 1st of 27 in the EU, every year checked. On "how do you expect your finances to change next year," also 1st of 27, every year checked. On general life satisfaction — a broader, less financially-loaded question — Greece

ranks somewhere between 2nd and 6th worst, never the most extreme. Put all three on the same standardized scale, so they're comparable rather than just ranked, and Greece's distance from the EU average on the two money questions comes out roughly three times larger than its distance on life satisfaction. A generic cultural pessimism doesn't pick and choose its topic. Whatever this is, it's aimed specifically at the financial questions, not at everything Greek households are asked. That rules out one explanation and not another. It rules out a general national gloom. It does not rule out a reporting difference that is itself about money — a sense, say, that the tax system and the crisis settlement fell unfairly — because that too would show up on the financial questions and nowhere else. This test cannot tell the two apart.

Outside research adds two things worth carrying forward honestly. The broader concern behind this whole chapter — that some cross-country differences in self-reported wellbeing really are a matter of response style rather than substance — is a documented, real phenomenon in the survey-methods literature, which is exactly why this question earned its own chapter instead of being waved off. And independent, Greece-specific research (a Crisis Observatory/ELIAMEP social-profile study; a 2015 academic account of Greece's welfare state before and after the crisis) backs up both halves of the finding here on its own terms: a real structural reason for the pre-crisis baseline — a thinner social safety net than most of the EU, more weight resting on families — and real material deterioration after 2009, not a survey artifact standing in for it.

So: a pre-existing baseline, real and conceded, sits alongside a crisis-era widening that is larger than any other EU country's, concentrated specifically on financial questions rather than spread across every measure of wellbeing, and independently corroborated by researchers who never touched this project's own data. Culture may set part of where Greece started. It does not explain where Greece has ended up, or how it got there — and the honest way to put it is not "Greeks are pessimists," but that most of what's left, once comparable conditions and accumulated exposure (Chapter 8) are already accounted for, isn't mood at all.

Read the method and caveats

METHOD NOTE

Six-test battery: coverage-audited pre-crisis ranking across three balanced panels (full-coverage six, 2005–2008, 2007–2008); a difference-in-differences model with an event-study check for pre-existing trends, alternative crisis-start dates (2009/2010/2011), a country-by-country placebo test across all 27 EU countries, leave-one-country-out stability, a periphery-only comparison group, and a synthetic-control check that did not work and is reported as such; plus a standardized (not just ranked) cross-indicator comparison against financial expectations and life satisfaction, and a specification-sensitivity check confirming Chapter 8's residual trend survives three different model specifications. Full derivation and every exact figure: `scripts/40–42_reporting_style_robustness*.py` in the technical report's pipeline, logged in full in its `publication_strategy.md`.

Landing

Greece is not a country that simply feels poorer than it is. The standard income-poverty measure, AROP, ranks it 4th-highest of 27 EU countries — elevated, not extreme. Subjective poverty ranks it 1st, by the widest margin in Europe, every year since 2011. Part of that 47.6-point gap is arithmetic: AROP's own threshold fell along with Greek incomes, missing the sharpest part of the collapse. The EU's broader official measure, AROPE, sees more of the damage — but still sits 39.7 points below what households report, the largest gap of any EU country either way. Even after correcting the AROP yardstick and testing the gap against comparable countries, a real residual survived, all the way through Chapter 7.

That remaining gap turned out not to be one hidden variable. It was a sequence: **the EU's deepest and longest unresolved pre-crisis GDP shortfall, a paycheck that fell further, ordinary prices made heavy by low pay, housing costs that reached even outright owners, joblessness that became long-term, and hardship that holding a job does not automatically resolve.** In the fixed model, sustained excess unemployment over roughly the past decade closes the remaining residual. In the stricter test, each fold chooses its own duration measure without seeing the country it must predict; Greece then lands at +2.70 points, 19th of 27, with eight countries predicted less accurately. The family of duration measures, not one hand-picked variable, is the durable result. Financial pessimism sits on top of that history — not the explanation, but a reasonably accurate memory of having lived through it.

What can look from outside like exaggerated Greek pessimism is better described as reported hardship consistent with prolonged material exposure. The fixed model slightly overpredicts Greece; the fully nested procedure leaves a small positive residual. Neither result is a verdict on national psychology. Together they show that Greece is no longer an exceptional reporting outlier once duration is counted.

There is one lesson here that reaches beyond Greece. AROP is a reasonable way to measure poverty relative to today's median income — but in a country living through a prolonged, broad-based collapse, a ruler that moves with the thing it's measuring can understate the damage for years at a time, and that is not a flaw unique to Greece or to this report. The practical takeaway is not "distrust AROP." It's that in crisis conditions, AROP is safest read alongside a fixed-standard measure, a broader official indicator like AROPE, and some account of how long the exposure has actually lasted — not on its own.

This is not a closed case. A small residual remains under the fully nested procedure; household composition, informal income, savings buffers, and institutional factors beyond Eurostat remain open. Everything here uses aggregate country-level data, not household microdata, and every relationship is associational. The defensible conclusion is narrower: Greece no longer looks exceptional once sustained exposure is counted, and the part the evidence explains is not generic mood.

historical data. Institutional-trust figures are cited context, not modeled inputs. Data vintage: 2026-08-19 for core series; 2026-08-20 for work effort.